

Assignment 1

Håkon Mork
ESCE 526 Artificial Intelligence

February 4, 2013

1

1.1 Number of states visited with simple heuristic

Game A:

Cutoff depth	3	4	5	6
Minmax	77,445	1,276,689	21,335,620	Unknown
α - β pruning	4129	48,203	694,652	Unknown
Improvement	$\times 18.76$	$\times 26.49$	$\times 30.72$	Unknown

Game B:

Cutoff depth	3	4	5	6
Minmax	98,345	1,704,319	29,770,996	Unknown
α - β pruning	6421	96,884	1,683,194	Unknown
Improvement	$\times 15.32$	$\times 17.59$	$\times 17.69$	Unknown

Game C:

Cutoff depth	3	4	5	6
Minmax	69,954	1,237,535	22,191,032	Unknown
α - β pruning	3763	51,098	840,633	Unknown
Improvement	$\times 18.59$	$\times 24.22$	$\times 26.40$	Unknown

This game has a high branching factor, in the ballpark of about 16 possible moves for each round, which makes exploring the game tree to any significant depth a very time-consuming chore. I didn't have time to calculate the numbers for cutoff depth 6, but they should be at least one order of magnitude higher than for the previous depth. Depressingly, I also found a bug in my program shortly before the submission deadline, so I'm worried that I have miscalculated the minmax figures. An interesting observation is that alpha-beta pruning seems to be increasingly efficient at chopping off irrelevant branches of the game tree as the cutoff

depth increases.¹ Observe, for example, that in game A the number of states visited by minmax with cutoff depth 3 was 18 times higher than the number visited by alpha-beta, while it was over 30 times higher when the cutoff depth was 5.

1.2 Does state generation order matter?

My evaluation function iterates through the successor states in the order they were generated: left-to-right, top-to-bottom, with the directions generated in the (arbitrary) order north-east-south-west. I considered the first move in game A and ran² alpha-beta pruning five times with a cutoff depth of 3, shuffling the list of successor states randomly every time one is generated. (Since the minmax algorithm does not prune the game tree at all, the order in which it evaluates successors is irrelevant.)

I found that evaluation order *did* matter, though not impressively so. Alpha-beta pruning with the non-shuffled evaluation order visited 4129 states, as per the table above. The sample runs with suffling visited 4480, 4338, 4324, 4114, and 4376 states, respectively, for an average of 4326 states, which is 4.77% more than the non-shuffled case. The maximum deviations from the original order were 8.5% more and 0.36% fewer states visited. (I also tried evaluating states in the reverse order; the difference was negligible.) This suggests that there is no significant benefit to be gained by shuffling; in fact, we see that more states were visited with alternative evaluation orders than with the original sequence. I don't know if this is a coincidence or if the order I picked is somehow optimal; alpha-beta pruning should perform best when the most valuable branches of the game tree are explored early, and most of the subsequent branches are pruned.

2

2.1 Choice of evaluation function

The fancyheuristic function calculates a heuristic value for a given player in a given board state. The heuristic has two parts: one to give a score based on how good the player's position is, and one to give a score for foiling the other player's plans. I underestimated the time it would take to implement a good heuristic, and due to coursework in other subject I didn't get time to work on it as much as I would have liked.

The part of the heuristic that evaluates the value of the board positions is rather simple. Let n_i be the number of i -in-a-row instances the player has on the board, either horizontally, vertically, or diagonally. For example, if the player has 3 pieces in a row at two different spots on the board, we have $n_3 = 2$. This number is multiplied by a corresponding power of 10, so that more weight is given to board states with more pieces connected. That is, the score is $\sum_{i=2}^4 n_i \cdot 10^i$. This ensures that the player gets a higher score for having more pieces in a row; that is, having several instances of two pieces connected is by far outweighed by having three pieces connected. Though this is hardly an optimal weighting, from my observations it looks like it performs reasonably well.

¹In my code I consider the children of the current state to be in ply 0; their children, i.e., the grandchildren of the current state, are in ply 1, and so on. If this interpretation is incorrect and the current state's children should instead be regarded as being on ply 1, the values in the tables above should be shifted one column to the right.

²The commands given were `python ass1.py --input starta.txt --cutoff 3 --alg ab --count --shuffle`. See the appendix for details on usage.




Figure 1: Example position.

The heuristic also comprises the sabotage function, which gives bonus points for obstructing the opponent. If the opponent has three pieces in a row and is just about to win, placing a piece such that an end of that line is blocked confers a big bonus on that game state’s heuristic value. The rationale is that only one thing should be better than hindering the opponent from connecting four pieces: namely, to connect four of one’s own pieces. Therefore the penalty or bonus is 9999, which is barely smaller than the 10,000 gained from winning. I haven’t been able to say conclusively whether the agent plays better with this addition, but it stands to reason that it should.

As an example of how the heuristic works, consider the situation in figure 1, where the white player gets $1 \times 10^3 + 2 \times 10^2 = 1200$ points for having that many pieces in a row. The black player gets a bonus for blocking a white connect-4 while collecting $1 \times 10^3 + 1 \times 100 = 1100$ points for the rest of the pieces.

In a sense, my agent follows a greedy algorithm: it follows a (locally) utility-maximizing sequence of actions. Thus, since winning the game has the highest possible heuristic value, the agent will try to achieve that outcome as soon as it can, without any specific regard to the inevitability (or not) of that outcome. Likewise, since losing has low utility but blocking the opponent from winning has high utility, the agent will try to prevent the opponent’s victory. It does not specifically try to delay its own defeat, other than in terms of choosing the course of action that maximizes utility. Hopefully those concerns will coincide.

2.2 Number of states visited with advanced heuristic

The improved heuristic does *not* consistently reduce the number of states visited, which I found surprising: for example, with alpha-beta pruning and a cutoff depth of 3, it visited 11,152 states, compared to the simple heuristic’s 4129 in the first move of game A. With a cutoff depth of 4 it visited 49,254, which is comparable to the simple heuristic’s 48,203. With minmax it visits the same number of states since the game tree isn’t pruned. While it may be the case that my improved heuristic just isn’t any good, it could also be that a significant number of states have the same heuristic value. I find it hard to believe that there should be more of those than all the states with value 0 in the case of the simple heuristic, so I have no plausible explanation.

2.3 Tradeoff between evaluation function and game tree depth

The improved heuristic is significantly more computationally intensive, and takes longer to explore the game tree to a given depth than the simple heuristic. An agent using the simple one can thus explore a greater number of states within a given amount of time. While this is certainly a drawback for the advanced heuristic, I found that the agent’s gameplay improved significantly, even when severely constrained by time or cutoff depth. In contrast, the simple heuristic, even when given more time and a greater depth, saw my agent explore a vastly greater game tree and still come up with obviously boneheaded decisions. Thus it seems to me that improving the quality of the heuristic is preferable to throwing more hardware at a simple, stupid heuristic to allow it to look even further ahead and explore the state space even more. That is, as long as we can’t have our cake and eat it too.

A Appendix: Source code

A.1 Implementation comments

The implementation is single-threaded, which is obviously not optimal in terms of pure performance. It seems to me that exploring a state space is a problem that lends itself well to parallelization since the subproblems are independent and potentially non-overlapping.

A.2 Usage

All arguments are optional:

- `-i` or `--input`: Specify an input file to be used as the initial game state. A plain-text file following the notation used in the assignment is expected. Defaults to the example illustrated in the “Introduction” part of the assignment text.
- `-u` or `--human`: The computer should play against a human adversary, not just against itself. May take values `w` or `b` to indicate that the human should be white or black, respectively.
- `-c` or `--cutoff`: Specify a cutoff depth. Defaults to 3.
- `-a` or `--alg`: Specify which of the minmax or alpha-beta pruning algorithms is to be used. May take values `mm` or `ab`. Defaults to alpha-beta pruning.
- `-l` or `--log`: A log file should be written on exit. May prove useful for the tournament.
- `-k` or `--count`: Count the number of states visited. Used for problem 1.1.
- `-s` or `--shuffle`: Shuffle the list of successor states before evaluating them. Used for problem 1.2.
- `-f` or `--fancy`: Indicate that the advanced heuristic should be used.
- `-t` or `--time`: Time limit (in seconds) for each move.

Example: `python ass1.py --input file.txt --alg ab --human w --time 20 --fancy --log`

A.3 Listing

The code is written in Python 2.7. I’ve expunged all logging statements and other debugging aids from this listing for the sake of readability. I find that Python is quite legible in most cases, so I’ve primarily added comments to explain purpose.

```
1 #!/usr/bin/env python
2 # -*- coding: utf-8 -*-
3 # Usage: python ass1.py [args]
4 # See submitted report for details.
5
6 import string, copy, time, argparse, random
```

```

7
8 # Tuples of (dy, dx) for all directions
9 directions = {
10     "N": (-1, 0),
11     "E": (0, 1),
12     "S": (1, 0),
13     "W": (0, -1)
14 }
15
16 # Used for counting states, problem 1.1
17 statesvisited = 0
18
19 # Used for the search algorithms.
20 class Node:
21     def __init__(self, board, player, command):
22         self.board = board
23         self.player = player
24         self.value = fancyheuristic(board, player) if fancy else
                simpleheuristic(board, player)
25         self.command = command # The move made to generate this state
26
27 # Dummy classes for representing the players.
28 class Black:
29     def __init__(self):
30         self.piece = "X"
31
32 class White:
33     def __init__(self):
34         self.piece = "O"
35
36 # Generates a list of all possible successor states to the given board
    position.
37 def successors(board, player):
38     succs = []
39     for y, line in enumerate(board):
40         for x, char in enumerate(line):
41             if char == player.piece:
42                 # Try all possible moves: xyN, xyE, xyS, xyW
43                 for cmd in (str(x + 1) + str(y + 1) + d for d in directions):
44                     try:
45                         candidate = move(cmd, board, player)

```

```

46         succs.append(Node(candidate, player, cmd))
47     except (ValueError, IndexError) as e:
48         # ValueError: attempted move was illegal, e.g. trying
49         # to move to an occupied square
50         # IndexError: try to move outside of the board
51         continue
52 # Used for problem 1.2, for determining whether varying the evaluation
53 # order matters
54 if args.shuffle:
55     random.shuffle(succs)
56 return succs
57
58 def alphabeta(player, node, depth, alpha, beta):
59     if countingstates:
60         global statesvisited
61         statesvisited += 1
62     succs = successors(node.board, player)
63     otherplayer = white if player is black else black
64     # Cut off and return heuristic value if we are too deep down
65     if depth == cutoff or len(succs) == 0:
66         return node.value
67     # White is maxplayer (arbitrary pick)
68     elif player is white:
69         for childnode in succs:
70             alpha = max(alpha, alphabeta(otherplayer, childnode, depth + 1,
71             alpha, beta))
72             if alpha >= beta:
73                 return beta
74         return alpha
75     # Black is minplayer
76     else:
77         for childnode in succs:
78             beta = min(beta, alphabeta(otherplayer, childnode, depth + 1,
79             alpha, beta))
80             if alpha >= beta:
81                 return alpha
82         return beta
83
84 def minmax(player, node, depth):
85     otherplayer = white if player is black else black
86     if countingstates:

```

```

83     global statesvisited
84     statesvisited += 1
85     if depth == cutoff or not successors(node.board, player):
86         return node.value
87     elif node.player is black: # Arbitrary pick
88         return min(minmax(otherplayer, child, depth + 1) for child in
91                     successors(node.board, player))
89     else:
90         return max(minmax(otherplayer, child, depth + 1) for child in
91                     successors(node.board, player))
92 # Returns a comma-separated board of X-es and O-s to be printed to console.
93 def prettyprint(board):
94     b = "\n".join(",".join(map(str, row)) for row in board)
95     return b.replace("None", " ")
96
97 # Check if any consecutive n entries in a row are X-es or O-s, and
98 # return the number of n-in-a-row instances on the board for the given player.
99 def horizontal(board, n, player):
100     piece = player.piece
101     connected = 0
102     for line in board:
103         for i, char in enumerate(line):
104             if line[i : i + n] == [piece] * n:
105                 connected += 1
106     return connected
107
108 # Checking verticals is equivalent to checking horizontals in the transposed
    matrix.
109 def vertical(board, n, player):
110     return horizontal(map(list, zip(*board)), n, player)
111
112 # All downward diagonals must start in the upper-left 4x4 submatrix, and
    similarly, all upward diagonals must start in the lower-left 4x4
    submatrix.
113 # Somewhat inelegant, but it works.
114 def diagonal(board, n, player):
115     piece = player.piece
116     connected = 0
117     for i in range(n):
118         for j in range(n):

```

```

119         # Count four down(up)ward from the upper (lower) left quadrant.
120         if (all(board[i + k][j + k] == piece for k in range(n))
121             or all(board[6 - i - k][j + k] == piece for k in range(n))):
122             connected += 1
123     return connected
124
125 # Indicate the winner (if any) in the given board state.
126 # Used, among other things, for the main game loop, which runs as long as
    there is no winner.
127 def winner(board):
128     if horizontal(board, 4, white) or vertical(board, 4, white) or
        diagonal(board, 4, white):
129         return white
130     elif horizontal(board, 4, black) or vertical(board, 4, black) or
        diagonal(board, 4, black):
131         return black
132     else:
133         return None
134
135 # Indicated whether the player has managed to thwart the opponent's play by
    blocking three of their pieces, thus preventing a loss.
136 # Used by the advanced utility function.
137 def sabotage(board, player):
138     goal = "000X" if player is black else "XXX0"
139     # This is a terrible, terrible hack, and I'm ashamed of it.
140     # Map the elements in the matrix to strings, concatenate,
141     auxboard = [map(str, l) for l in copy.deepcopy(board)]
142     auxboard = ["".join(l) for l in auxboard]
143     auxtransp = ["".join(l) for l in map(list, zip(*auxboard))]
144     # then look up XXX0 and 000X and their reverses in that string.
145     hor = any(goal in line or goal[::-1] in line for line in auxboard)
146     vert = any(goal in line or goal[::-1] in line for line in auxtransp)
147     # The diagonal is a bit more tricky, but the same reasoning applies as in
        the horizontal(board, n, player) function.
148     # All interesting diagonals start in the upper or lower left quadrants,
        so we make a list of them, join each of them up and look for the 000X
        and XXX0 strings and their reverses there.
149     diags = []
150     for i in range(4):
151         for j in range(4):
152             diags.append([board[i + k][j + k] for k in range(4)])

```



```

153         diags.append([board[6 - i - k][j + k] for k in range(4)])
154     # Map elements to string and concatenate with empty string.
155     diags = ["".join(l) for l in [map(str, l) for l in diags]]
156     diag = any(goal in line or goal[::-1] in line for line in diags)
157     return hor or vert or diag
158
159 # As given in problem 1.
160 def simpleheuristic(board, player):
161     otherplayer = white if player is black else black
162     if winner(board) is player:
163         return 1
164     elif winner(board) is otherplayer:
165         return -1
166     else:
167         return 0
168
169 # A somewhat more advanced heuristic, used for part 2 of the assignment and
170 # actual gameplay. See the submitted report for details and discussion.
171 def fancyheuristic(board, player):
172     otherplayer = white if player is black else black
173     score = 0
174     for i in [4, 3, 2]:
175         h = horizontal(board, i, player)
176         v = vertical(board, i, player)
177         d = diagonal(board, i, player)
178         score += (10 ** i) * (h + v + d)
179     if sabotage(board, player):
180         score += 9999
181     elif sabotage(board, otherplayer):
182         score -= 9999
183     return score
184
185 # Builds a matrix from an input string, in case we want to specify an initial
186 # board layout.
187 def parseboard(boardstring):
188     boardstring = string.replace(boardstring, ",", "")
189     board, line = [], []
190     for char in boardstring:
191         if char == " ":
192             line.append(None)
193         elif char == "\n":

```

```

192         board.append(line)
193         line = []
194     else:
195         line.append(char)
196     if line:
197         board.append(line) # Last line, if there is no newline at the end
198     return board
199
200 # Performs a move according to a given command, and returns the new game
    state.
201 def move(command, board, player):
202     # Takes indices and a direction, e.g. "43W" or "26N".
203     x, y, d = tuple(command)
204     # The board is a zero-indexed array, adjust accordingly
205     x, y = int(x) - 1, int(y) - 1
206     dy, dx = directions[d.upper()]
207     # Does the piece fall within the bounds?
208     if ((0 <= x + dx <= 7) and (0 <= y + dy <= 7)
209         # ...and is it our piece?
210         and board[y][x] == player.piece
211         # ...and is the destination square empty?
212         and not board[y + dy][x + dx]):
213         # ...then it's okay
214         successor = copy.deepcopy(board)
215         successor[y + dy][x + dx] = successor[y][x]
216         successor[y][x] = None
217         return successor
218     else:
219         raise ValueError("The move " + command + " by " +
            player.__class__.__name__ + " is not legal")
220
221 # Parse command-line arguments. See submitted report for summary of usage.
222 parser = argparse.ArgumentParser()
223 parser.add_argument("-c", "--cutoff", help="Cutoff depth")
224 parser.add_argument("-i", "--input", help="Input game board")
225 parser.add_argument("-u", "--human", choices=["w", "b"], help="Play with a
    human opponent")
226 parser.add_argument("-a", "--alg", choices=["mm", "ab"], help="Minmax or
    alpha-beta algorithm")
227 parser.add_argument("-l", "--log", help="Write a game log on exit",
    action="store_true")

```

```

228 parser.add_argument("-s", "--shuffle", help="Shuffle successor list",
    action="store_true")
229 parser.add_argument("-k", "--count", help="Count states visited",
    action="store_true")
230 parser.add_argument("-f", "--fancy", help="Fancy heuristic function",
    action="store_true")
231 parser.add_argument("-t", "--time", help="Timeout limit in seconds")
232 args = parser.parse_args()
233
234 cutoff = int(args.cutoff) if args.cutoff else 3
235 useab = not (args.alg == "mm") # Alpha-beta by default
236 logthegame = args.log
237 countingstates = args.count
238 fancy = args.fancy
239 timeout = float(args.time) if args.time else float("inf")
240
241 # If we give an input file, parse it and set up the initial board layout
    accordingly.
242 if args.input:
243     with open(args.input, "r") as inputfile:
244         initstr = inputfile.read()
245         board = parseboard(initstr)
246 else: # If not, default to the starting position from the assignment.
247     board = [
248         ["O", None, None, None, None, None, "X"],
249         ["X", None, None, None, None, None, "O"],
250         ["O", None, None, None, None, None, "X"],
251         ["X", None, None, None, None, None, "O"],
252         ["O", None, None, None, None, None, "X"],
253         ["X", None, None, None, None, None, "O"],
254         ["O", None, None, None, None, None, "X"]
255     ]
256
257 # Player instances
258 white = White()
259 black = Black()
260
261 # Designate a human player to a color if one is given, else let the computer
    play against itself.
262 if args.human == "w":
263     human = white

```

```

264     computer = black
265 elif args.human == "b":
266     human = black
267     computer = white
268 else:
269     human = None
270     computer = black # Arbitrary choice
271
272 # Other administrivia
273 currentplayer = white
274 log = ["Initial state:"]
275 movenumber = 1
276
277 # Main loop. Runs as long as there is no winner, or until interrupted.
278 while winner(board) is None:
279     # Print informative stuff at the beginning of each round.
280     playername = currentplayer.__class__.__name__
281     p = prettyprint(board)
282     print p
283     print "\nMove #s:" % movenumber
284     print "It's %s's turn." % playername
285     if logthegame:
286         log.append(p)
287         log.append("\nMove #s:" % movenumber)
288         log.append("It's %s's turn." % playername)
289
290     cmd = "" # Command string, e.g. 11E or 54N
291
292     try: # In case of keyboard interrupts
293         # Show a list of options to human players:
294         if currentplayer is human:
295             print "Possible moves:"
296             for s in successors(board, currentplayer):
297                 print s.command
298             cmd = raw_input()
299         # Otherwise, have the computer calculate its move using the given
300         # algorithm.
301     else:
302         t = time.time() # Time limit is 20 seconds
303         succs = successors(board, currentplayer) # Successors of this
304             state

```

```

303         # Take the first move, pick something better later on if we can
           find it.
304     bestmove = succs[0].command
305     bestutility = 0
306
307     # Pick algorithm according to --alg argument.
308     if useab: # Alpha-beta pruning
309         for succ in succs:
310             # Initialize with alpha = -inf, beta = inf
311             u = alphabeta(currentplayer, succ, 0, float("-inf"),
                           float("inf"))
312             if u > bestutility:
313                 bestutility = u
314                 bestmove = succ.command
315             if time.time() - t > timeout:
316                 print "Time limit cutoff"
317                 break
318     else: # Minmax
319         for succ in succs:
320             u = minmax(currentplayer, succ, 0)
321             if u > bestutility:
322                 bestutility = u
323                 bestmove = succ.command
324             if time.time() - t > timeout:
325                 print "Time limit cutoff"
326                 break
327
328     cmd = bestmove
329     print "The computer makes the move", cmd
330     print "Thinking took", time.time() - t, "seconds"
331     if logthegame:
332         log.append("Thinking took " + str(time.time() - t) + "
           seconds")
333
334     # May raise a ValueError if input is ill-formed.
335     board = move(cmd, board, currentplayer)
336
337     if countingstates:
338         print statesvisited
339         raise Exception("Counting states only, stopping here")
340     if logthegame:

```

```

341         log.append("%s plays %s" % (playername, cmd))
342
343         # Move to next round
344         currentplayer = white if currentplayer is black else black
345         playername = currentplayer.__class__.__name__
346         movenumber += 1
347
348         # Catch errors made by user when entering a command:
349         except ValueError:
350             print "Illegal move."
351         # Possibility for interrupting computation it takes too long.
352         except KeyboardInterrupt:
353             if logthegame:
354                 log.append("Game cancelled.")
355             break
356
357 # Post-game formalities: print the board one last time, logging
358 print prettyprint(board)
359
360 if winner(board):
361     s = "%s won the match" % winner(board).__class__.__name__
362     print s
363     if logthegame:
364         log.append(s)
365 else:
366     print "It's a draw"
367     if logthegame:
368         log.append("It's a draw")
369
370 if logthegame:
371     log.append(prettyprint(board))
372     logname = time.strftime("./connect4-%H-%M-%S.log")
373     with open(logname, "w+") as logfile:
374         logfile.write("\n".join(log))

```